

Naïve Bayes Model

Question 1:

Language	Prior	Count	Count/Total Count
TEL	0.099329109	533	0.099329109
HIN	0.065598211	352	0.065598211
SPA	0.083861349	450	0.083861349
KOR	0.103801714	557	0.103801714
FRA	0.088147596	473	0.088147596
JPN	0.103801714	557	0.103801714
ARA	0.092061126	494	0.092061126
ITA	0.096161014	516	0.096161014
TUR	0.093924711	504	0.093924711
ZHO	0.110510622	593	0.110510622
DEU	0.062802833	337	0.062802833

Question 2:

Label	priorProbs	Label Count
TUR	0.095317726	57
JPN	0.100334448	60
ARA	0.085284281	51
DEU	0.056856187	34
KOR	0.100334448	60
TEL	0.10367893	62
HIN	0.078595318	47
SPA	0.086956522	52
ITA	0.088628763	53
ZHO	0.115384615	69
FRA	0.088628763	53

If we classified doc in the dev by the classifier with the highest prior probability from the training data set (ZHO with $P(Y) = 0.110510622$), the majority class baseline accuracy would be 0.115384615.

Question 3:

alpha	accuracy
0.01	0.725752508
0.05	0.735785953
0.1	0.747491639
0.2	0.74916388
0.5	0.732441472
1	0.68729097
2	0.581939799
5	0.382943144

Based on the alpha tuning, $\alpha = 0.2$ gives the best accuracy.

Question 4:

alpha	accuracy
0.01	0.719063545
0.05	0.737458194
0.1	0.737458194
0.2	0.732441472
0.5	0.7090301
1	0.678929766
2	0.575250836
5	0.377926421

Lemmatization lowers the optimal α by almost 2%. A possible reason is that certain translations of words are used by different native speakers.

Question 5:

dev: 0.7491638795986622
test: 0.7152317880794702

Using the $\alpha = 0.05$ from lemmatization
dev: 0.7374581939799331
test: 0.6771523178807947